

# **ENHANCING THE PERFORMANCE OF E-COMMERCE BUSINESSES**

This report explores key strategies for e-commerce success: improving customer retention, increasing product revenues, and optimizing inventory management. It leverages extensive datasets from platforms like Shiprocket and INCREFF, covering financial metrics, SKU codes, inventory levels, and product details from major stores such as Ajio, Amazon, Flipkart, and others. Analyses focus on customer purchases, transaction rates, categories, fulfilment methods, and gross amounts.

Three main datasets—"International\_sales\_Report.csv," "Amazon\_Sale\_Report.csv," and "Sales\_Report.csv"—provide insights into sales quantity, pricing, order details, and product performance. Identified areas for further investigation are accompanied by concise problem statements for targeted analysis.

## **Problem Statement 1 - Comparing the profitability of one-time VS repeat customers**

Analysing customer retention rates is crucial for businesses aiming to improve loyalty which is vital for sustained profitability and long-term success. It measures the percentage of customers who are consistent with purchases. Strategies to boost customer satisfaction and loyalty can be developed by studying past purchases history and consequently, company discern patterns for the future. This involves examining preferences, purchase frequency, service interactions, and other factors influencing repeat business.

### **Data Cleaning & Manipulation**

I had to thoroughly clean the dataset before analysing its structure and contents. I noticed that rows 18637 to 19676 were empty, and the columns were in the wrong order. To tackle this, I split the dataset and planned to rejoin it later during cleaning. Additionally, I used the slice() function to remove rows containing irrelevant information, as shown in Figure 1.

```
> correct_part <- int_sale_rep_data %>% slice(1:18635)

| > wrong_part <- int_sale_rep_data %>% slice(19678:nrow(int_sale_rep_data)-1)
```

*Figure 1*

I rearranged the columns correctly using the data.frame() function, creating a new data frame to hold the required information. The size column was filled by extracting the last digits from the SKU codes using regular expressions (RegEx) and the grepl() function to search

for matches in the string vector. Then, we rejoined the data frame with the previously sliced dataset, as illustrated in Figure 2.

```
> wrong_part_corrected <- data.frame(index = wrong_part$index, DATE = wrong_part$Months, Months = wrong_part$CUSTOMER, CUSTOMER = wrong_part$DATE, Style = wrong_part$Style, SKU = wrong_part$SKU, Size = NA, PCS = wrong_part$Size, RATE = wrong_part$PCS, GROSS.AMT = wrong_part$RATE)

> wrong_part_corrected <- wrong_part_corrected %>% mutate(Size = ifelse(grepl("-", SKU), sub(".*-([A-Za-z]+)$", "\\1", SKU), ""))
>
> # Rejoin the dataframe
> int_sale_rep_data <- bind_rows(correct_part, wrong_part_corrected)
```

Figure 2

After data cleaning, the total observations in this dataset is 34598 and the total number of variables is 10. The international sales data contains 9 categorical variables and 1 numerical variable based on the datatypes shown in Figure 3.

```
> str(int_sale_rep_data)
'data.frame':   34598 obs. of  10 variables:
 $ index      : int   0 1 2 3 4 5 6 7 8 9 ...
 $ DATE       : chr   "06-05-21" "06-05-21" "06-05-21" "06-05-21" ...
 $ Months     : chr   "Jun-21" "Jun-21" "Jun-21" "Jun-21" ...
 $ CUSTOMER   : chr   "REVATHY LOGANATHAN" "REVATHY LOGANATHAN" "REVATHY LOGANATHAN" "REVATHY L
OGANATHAN" ...
 $ Style      : chr   "MEN5004" "MEN5004" "MEN5004" "MEN5009" ...
 $ SKU        : chr   "MEN5004-KR-L" "MEN5004-KR-XL" "MEN5004-KR-XXL" "MEN5009-KR-L" ...
 $ Size       : chr   "L" "XL" "XXL" "L" ...
 $ PCS        : chr   "1.00" "1.00" "1.00" "1.00" ...
 $ RATE       : chr   "616.56" "616.56" "616.56" "616.56" ...
 $ GROSS.AMT  : chr   "617.00" "617.00" "617.00" "617.00" ...
```

Figure 3

1 variable contains missing values (NAs), which is the SKU variable. The missing value count is shown in Figure 4 below. Empty spaces were also replaced with NA to ensure it was calculated and removed. The missing values in the dataset were removed using the `na.omit()` function. I eliminate any missing values initially identified in the dataset to prevent errors, anomalies, or biases that may arise from their absence. This process also helps reduce the datasets complexity.

```
> # Check for NA values and remove them
> int_sale_rep_data[int_sale_rep_data == ""] <- NA
> colSums(is.na(int_sale_rep_data))
index      DATE      Months  CUSTOMER      Style      SKU      Size      PCS
0          0          0          0          0      1320          0          0
RATE GROSS.AMT
0          0

>
> int_sale_rep_data <- na.omit(int_sale_rep_data)
```

Figure 4

The summary and descriptive statistics of the variables in cleaned data are shown in Figure 5 below. There are now 33278 observations after removing the missing values.

```

> # Summary statistics of cleaned data
> summary(international_cleaned)

```

DATE		Months	CUSTOMER	Style
Min.	:2021-06-05	Length:33278	Length:33278	Length:33278
1st Qu.	:2021-09-18	Class :character	Class :character	Class :character
Median	:2021-11-16	Mode :character	Mode :character	Mode :character
Mean	:2021-11-21			
3rd Qu.	:2022-02-21			
Max.	:2022-03-31			

SKU	Size	PCS	RATE
Length:33278	Length:33278	Min. : 1.000	Length:33278
Class :character	Class :character	1st Qu.: 1.000	Class :character
Mode :character	Mode :character	Median : 1.000	Mode :character
		Mean : 1.272	
		3rd Qu.: 1.000	
		Max. :15.000	

GROSS.AMT
Min. : 227.0
1st Qu.: 462.0
Median : 653.0
Mean : 829.4
3rd Qu.: 950.0
Max. : 9735.0

Figure 5

I performed a series of steps to manipulate the data for us to be able to analyse and compare the profitability of one – time and repeating customers. Firstly, I determined the date range of the sales in the dataset using the max() and min() functions – sales ranged from 5th June 2021 to 31st March 2022 as seen below.

```

> # Problem Statement 1: Comparing the profitability of one-time vs repeat customers
>
> # International sale report consist of sales from 5 June 2021 to 31 March 2022
> print(max(international_cleaned$DATE))
[1] "2022-03-31"
> print(min(international_cleaned$DATE))
[1] "2021-06-05"

```

Figure 6

As seen in Figure 7, the group\_by() function was then used to group customers and their purchase histories as I specify new columns through the use of the summarise() function.

```

> # Identify unique customers and their purchase histories
> customer_purchases <- international_cleaned %>%
+   group_by(CUSTOMER) %>%
+   summarise(
+     First_Purchase = min(DATE),
+     Last_Purchase = max(DATE),
+     Days_Since_Last_Purchase = max(DATE) - min(DATE),
+     Purchase_Count = n_distinct(DATE), # How many purchases made on different dates
+     Total_purchase_frequency = n(), # Total number of items purchased
+     Total_amount_spent = sum(GROSS.AMT) # Total amount spent (INR) per customer
+   )

```

Figure 7

Using the nrow() function, we know that there are 123 unique customers in the dataset. The filter() function is used to determine repeating customers where their purchase count was greater than 1 like seen below.

```

> total_customers <- nrow(customer_purchases) # 123 unique customers
> print(paste("Total unique customers:", total_customers))
[1] "Total unique customers: 123"
>
> repeat_customers <- customer_purchases %>% filter(Purchase_Count > 1)

> total_repeat_customers <- nrow(repeat_customers)
> print(paste("Total repeat customers:", total_repeat_customers))
[1] "Total repeat customers: 41"

```

Figure 8

Therefore, there are 41 repeating customers as seen in the output of the `nrow()` function used on the `repeat_customers` data frame. Subsequently, there are a total of 82 one-time customers as shown in the figure below.

```

> # 82 one time customers
> total_one_time_customers <- nrow(one_time_customers)
> print(paste("Total one-time customers:", total_one_time_customers))
[1] "Total one-time customers: 82"

```

Figure 9

Furthermore, I analysed the dataset by computing the average quantity per purchase for both one – time and repeating customers. As in Figure 10, this was performed through the use of `filter()` function for the specific type of customers as well as the `mean()` function within the `summarise()` function to output the required average value.

```

> # Print the average quantity per purchase for one time customers
> one_time_customers_avg_qty <- international_cleaned %>%
+   filter(CUSTOMER %in% one_time_customers$CUSTOMER) %>%
+   summarise(
+     Avg_Qty_Per_Purchase = mean(PCS)
+   )
> print("One-time Customers Average Quantity Per Purchase")
[1] "One-time Customers Average Quantity Per Purchase"
> print(one_time_customers_avg_qty) # 1.24
  Avg_Qty_Per_Purchase
1             1.238345

> # Calculate average quantity per purchase for repeat customers
> repeat_customers_avg_qty <- international_cleaned %>%
+   filter(CUSTOMER %in% repeat_customers$CUSTOMER) %>%
+   group_by(CUSTOMER) %>%
+   summarise(
+     Avg_Qty_Per_Purchase = mean(PCS)
+   )
> print("Repeat Customers Average Quantity Per Purchase")
[1] "Repeat Customers Average Quantity Per Purchase"
> print(repeat_customers_avg_qty %>% summarise(Avg_Qty_Per_Purchase = mean(Avg_Qty_Per_Purchase))) # 1.26
# A tibble: 1 x 1
  Avg_Qty_Per_Purchase
  <dbl>
1             1.26

```

Figure 10

Thus, it is clear that the repeating customers have a higher average quantity per purchase which denotes that more products are bought by repeating customers. To ensure the analysis is impactful, I further computed the summary of the total amount spent for both one – time and repeating customers, shown in the figure below. This was done using the total amount spent column in the dataset.

```

> # Do repeat customer really spend more than one-time customers?
> # Summary statistics of the total amount spent by one-time customers and repeat customers
> one_time_summary <- one_time_customers %>% select(Total_amount_spent) %>%
+ summary(Total_amount_spent)
> print("One-time Customers Spending Summary")
[1] "One-time Customers Spending Summary"
> print(one_time_summary)
Total_amount_spent
Min. : 10116
1st Qu.: 34722
Median : 60145
Mean : 145298
3rd Qu.: 133384
Max. : 885096

> repeat_summary <- repeat_customers %>% select(Total_amount_spent) %>%
+ summary(Total_amount_spent)
> print("Repeat Customers Spending Summary")
[1] "Repeat Customers Spending Summary"
> print(repeat_summary)
Total_amount_spent
Min. : 35370
1st Qu.: 119200
Median : 216686
Mean : 382569
3rd Qu.: 353076
Max. : 3325150

```

Figure 11

Hence, it is evident that repeat customers spend more money than the one-time customers as outputted in the summary of both customer categories. I notice a major difference in mean values of both spending summaries for one-time and repeat customers. I make this analysis clearer by creating a data frame which only holds information on the percentage of the total amount spent for both customer categories, seen below.

```

> total_spent <- data.frame(Group = c("One Time Customer", "Repeat Customer"),
+ count= c(sum(one_time_customers$Total_amount_spent),
+ sum(repeat_customers$Total_amount_spent)))
> total_spent <- total_spent %>%
+ mutate(percent = paste(100*round(count/sum(count), 3), "%"))
> head(total_spent)
  Group      count percent
1 One Time Customer 11914423 43.2 %
2 Repeat Customer 15685339 56.8 %

```

Figure 12

The repeating customer group spends a higher percentage compared to one-time customers, as seen in previous code. Additionally, I noticed that customers tend to return during specific times of the year. Using the `group_by()` function to group the data by sale date, it is evident that most customers return in January and around October, as depicted in Figure 13.

```

> # Analyze the distribution of repeat customers over time
> repeat_customer_trends <- repeat_customers %>%
+ mutate(YearMonth = floor_date>Last_Purchase, "month")) %>%
+ group_by(YearMonth) %>%
+ summarise(Repeat_Customer_Count = n()) %>%
+ arrange(desc(Repeat_Customer_Count))
>
> # Most of the customers repeated come back during January, a huge amount of customer also comes back around October
> # January for new year? October for Deepavali? These spike in increase repeat customer count may be caused by special occasions

```

Figure 13

Peak sales align with New Year and Deepavali celebrations, especially for Indian clothing like sarees and kurtas. Deepavali sales spiking in October confirms this trend. With 66.7% of customers being non-repeat purchasers, there's ample room for improvement. Repeat

customers, averaging 1.26 purchases, present an opportunity for bulk sales. To encourage repeat business, e-commerce firms should focus on improving customer satisfaction, potentially through loyalty programs and future purchase discounts.

### *Data Visualisations*

---

```
repeat_customers$metric <- repeat_customers$Total_amount_spent / repeat_customers$Total_purchase_frequency
one_time_customers$metric <- one_time_customers$Total_amount_spent / one_time_customers$Total_purchase_frequency
# Step 3: Calculate the average of the metrics
average_repeat <- mean(repeat_customers$metric, na.rm = TRUE)
average_one_time <- mean(one_time_customers$metric, na.rm = TRUE)

# Step 4: Create a data frame for plotting
average_data <- data.frame(
  Customer_Type = c("Repeat Customers", "One-time Customers"),
  Average_Metric = c(average_repeat, average_one_time)
```

*Figure 14*


The metric for repeat and one-time customers is calculated by dividing their total spent by their purchase frequency, stored in new columns. The average metric for each customer category is then computed using the `mean()` function with `na.rm = TRUE` to ensure accurate results by ignoring any missing values, enhancing analysis accuracy.

```
ggplot(average_data, aes(x = Customer_Type, y = Average_Metric, fill = Customer_Type)) +
  geom_bar(stat = "identity", width = 0.5) +
  labs(title = "Average Spending Per Purchase Based on Customer Types",
       x = "",
       y = "Average Spending Per Purchase") +
  theme_minimal() +
  scale_fill_manual(values = c("Repeat Customers" = "lightblue", "One-time Customers" = "pink")) +
  theme(legend.position = "none",
        plot.title = element_text(hjust = 0.5)) +
  geom_text(aes(label=round(Average_Metric, 2)),
            position=position_dodge(width=0.9), vjust=-0.25, size = 3)
```

*Figure 15*

A bar chart effectively illustrates the comparison of average spending between two customer categories by visually representing the difference in bar lengths. In Figure 15, I utilized the `ggplot` library to create a bar chart using `average_data` as the data source. Aesthetic mappings were set with `x` representing customer types, `y` for average metrics, and `fill` to differentiate bars by customer type. Adding bars to the plot with `geom_bar()`, I used the identity statistic to directly plot metric values and specified a width of 0.5. Additionally, I displayed average metric values above the bars using `geom_text()`, formatted to two decimal places and positioned appropriately.





*Figure 16*

Figure 16 illustrates the difference in average spending between one-time and repeat customers, showing that repeat customers tend to spend more on average per purchase; thus, repeat customers are likely more profitable for the company. Consequently, the company's priority should be increasing repeat customers, using strategies such as offering customer loyalty programs which could efficiently boost overall revenue, enhance profitability and build a more stable customer base.

A data frame is created like in Figure 17, where it contains two groups representing each customer type. Each group has its records of the sum for total\_amount\_spent from the initial dataset. The mutate function is used to add a new column percent to the total\_spent data frame. This column represents the percentage each group contributes to the total amount spent, calculated by dividing the count of each group by the total count and then formatted as a percentage.

```
total_spent <- data.frame(Group = c("One Time Customer", "Repeat Customer"),
                          count = c(sum(one_time_customers$Total_amount_spent),
                                     sum(repeat_customers$Total_amount_spent)))
total_spent <- total_spent %>%
  mutate(percent = paste(100*round(count/sum(count), 3), "%"))
```

*Figure 17*

A pie chart clearly shows the relative spending between the two. With only two categories, a pie chart with two slices is very simple and easy to understand at a glance. There is no risk of clutter, and the viewer can quickly grasp the information being presented. Hence, it effectively highlights the difference in spending between the two categories, making it visually apparent how the expenditures compare.

```
total_spent <- data.frame(Group = c("One Time Customer", "Repeat Customer"),
  count= c(sum(one_time_customers$Total_amount_spent),
    sum(repeat_customers$Total_amount_spent)))

total_spent <- total_spent %>%
  mutate(percent = paste(100*round(count/sum(count), 3), "%"))
View(total_spent)

ggplot(total_spent, aes(x="", y=count, fill=Group)) +
  geom_col(color = "black") +
  geom_text(aes(label = percent),
    position = position_stack(vjust = 0.5),
    size = 8,
    color = c("white", "black")) +
  coord_polar("y", start=0) +
  labs(title = "Total Spent by both Groups") +
  scale_fill_manual(values = c("pink", "lightblue")) +
  theme(plot.title = element_text(hjust = 0.5, size = 15),
    panel.border = element_blank(),
    axis.text = element_blank(),
    panel.grid = element_blank(),
    axis.ticks = element_blank(),
    panel.background = element_rect(fill = "#ebf2ff"),
    plot.background = element_rect(fill = "#ebf2ff"),
    legend.background = element_rect(fill = "#ebf2ff"),
    axis.title.y = element_blank(),
    axis.title.x = element_blank())
```

Figure 18

I utilized the ggplot library to generate a pie chart, as depicted in Figure 18. Initializing the plot with the total\_spent data frame, I defined aesthetic mappings where y=count represented bar heights and fill=Group distinguished colours. Adding coloured bars with black borders to the plot using geom\_col(), I included percentage labels on each bar positioned at the middle with geom\_text(). Transforming the plot into a circular polar coordinate system with coord\_polar("y", start=0), I labelled the plot with a title "Total Spent by both Groups" using labs(). Manually setting colours for the bars representing each group with scale\_fill\_manual(), I then adjusted various theme elements with theme(), such as removing the background grid and axis titles, and setting the plot background colour, to enhance appearance.

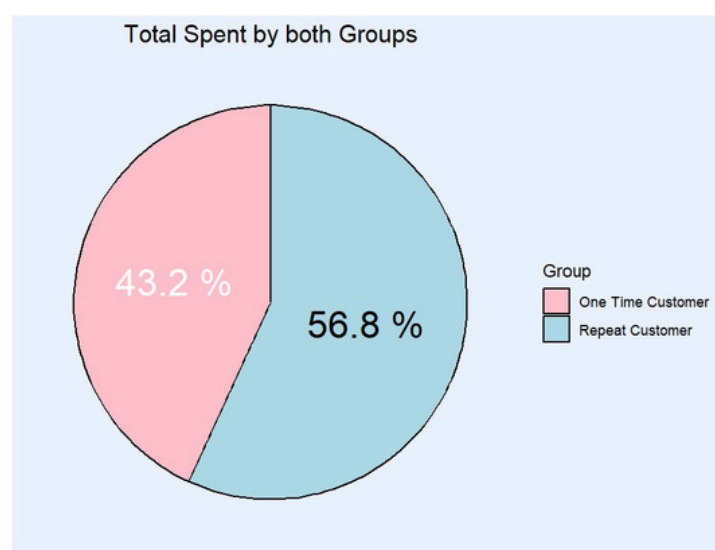


Figure 19



The pie chart illustrates total expenditures from two customer demographics: one-time and repeat customers. Repeat customers contribute a larger portion to the total expenditure; which is a significant portion of the business's revenue. Companies are able to analysis this data to comprehend its revenue sources and enhancing customer retention to foster repeat business. However, it's worth noting that while repeat customers spend more, the disparity in total spending between the two groups is relatively small, suggesting that their contribution is not significant compared to one-time customers.

Using the `repeat_customers` data frame, a new column called `YearMonth` is created with `mutate()`, rounding down `Last_Purchase` dates to the nearest month using `floor_date()` from the `lubridate` package. This effectively groups all dates within the same month. The data is then grouped by `YearMonth` with `group_by()`, and `summarise()` calculates the number of repeat customers for each month. `arrange()` sorts the data frame in descending order based on `repeat_customer_count`, showcasing months with the highest repeat customers first.

```
repeat_customer_trends <- repeat_customers %>%
  mutate(YearMonth = floor_date>Last_Purchase, "month")) %>%
  group_by(YearMonth) %>%
  summarise(Repeat_Customer_Count = n()) %>%
  arrange(desc(Repeat_Customer_Count))
```

*Figure 20*

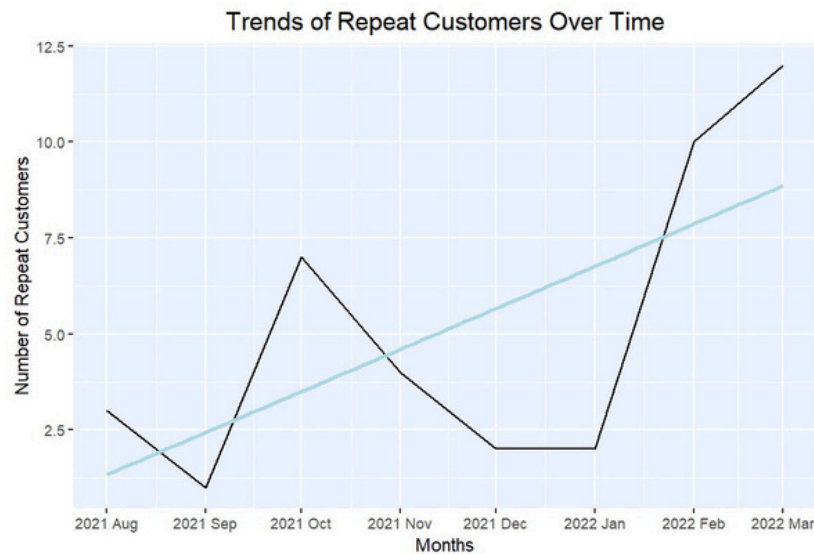
Line charts are particularly effective for illustrating trends over a period, making it easy to observe how data points change over intervals. They help in recognizing patterns or behaviours in data, such as seasonal fluctuations or long-term growth.

```
ggplot(repeat_customer_trends, aes(x = YearMonth, y = Repeat_Customer_Count)) +
  geom_line() +
  labs(title = "Trends of Repeat Customers Over Time", x = "Months", y = "Number of Repeat Customers") +
  scale_x_date(date_breaks = "1 month", date_labels = "%Y %b") +
  theme(plot.title = element_text(hjust = 0.5, size = 15),
        panel.background = element_rect(fill = "#ebf2ff")) +
  geom_smooth(method=lm, se=FALSE, col='lightblue', size=1)
```

*Figure 21*

I created a line chart using `ggplot`, with `repeat_customer_trends` data frame plotted against `YearMonth` on the x-axis and `Repeat_Customer_Count` on the y-axis. Adding a line plot layer with `Geom_line()`, I depicted trends over time. Customizing plot labels with `labs()`, I labelled the title as "Trends of Repeat Customers Over Time", x-axis as "Months", and y-axis as "Number of Repeat Customers". Formatting the x-axis with `scale_x_date()`, I set breaks to occur every month and labels to display year and month. Adjusting plot appearance, I centred the title, set its size, and modified the background colour of the

plotting area with theme(). Finally, I smoothed the line using linear regression without displaying standard error bands and customized the line colour and thickness with geom\_smooth().



*Figure 22*

The line graph illustrates the frequency of repeat customers over time, with notable peaks observed in October and February. These spikes coincide with significant holidays such as Deepavali in October and Lunar New Year in February. Leveraging these seasonal fluctuations could offer a strategic opportunity for the business to enhance both customer retention and satisfaction. By prioritizing efforts during these high-traffic periods, the business can capitalize on increased customer engagement and potentially further solidify its market presence.

### **Problem Statement 2 - Analysis of revenue made in each product category**

Discrepancies between shipment volume and revenue can signal inefficiencies or untapped opportunities. Understanding these gaps is vital for optimizing resource allocation and maximizing profitability. Factors like demand fluctuations, customer preferences, and average order values contribute to these differences. Comprehensive shipping data, including shipment numbers, product categories, destinations, and segmented revenue data, are essential for analysis. This problem statement utilizes the amazon\_sale\_report.csv dataset.

### Data Cleaning & Manipulation

Before I could delve into analysing the dataset, thorough data cleaning was imperative. Initially, I used the `is.na()` function to determine the number of missing values in the columns of the dataset

However, upon further checking through the dataset, I noticed that many data were depicted through empty strings thus not showcased in the count of missing values. Hence, it is required to replace the empty strings with NA values to be detected using the `is.na()` function.

```
> # Replace empty string values with NA
> amazon_sale_rep_data[amazon_sale_rep_data == ''] <- NA
>
> # Now we can obtain the real missing values
> colSums(is.na(amazon_sale_rep_data))
      index      Order.ID      Date
      0         0         0
      Status      Fulfilment      Sales.Channel
      0         0         0
ship.service.level      Style      SKU
      0         0         0
      Category      Size      ASIN
      0         0         0
      Courier.Status      Qty      currency
      6872         0      7795
      Amount      ship.city      ship.state
      7795         33         33
ship.postal.code      ship.country      promotion.ids
      33         33      49153
      828      fulfilled.by      Unnamed..22
      0      89698      49050
```

Figure 23

After analysing the dataset, I found 9 variables with numerous missing values. To further clean the dataset, I removed redundant columns like ship country and currency details. These columns didn't contribute to the analysis as information like ship\_city and ship\_state already implied the dataset was for India, making other columns irrelevant.

```
> # Remove redundant columns that are idea just for references or already implied
> # (ship country and currency we already know it's a dataset for india)
> amazon_sale_rep_data <- select(amazon_sale_rep_data, -index, -promotion.ids, -Unnamed..22, -fulfilled.by, -ship.country, -currency)
```

Figure 24

The dataset included information on the date however it was not in the accurate data type. Thus, conversion was done to the Date column using the `mutate()` and `as.Date()` function.

```
> # Convert Date to Date datatype
> amazon_sale_rep_data <- amazon_sale_rep_data %>%
+   mutate(Date = as.Date(Date, format="%m-%d-%y"))
```

Figure 25

After significant amount of data cleaning, the dataset consists of 128975 observations with 18 variables. There are 14 categorical variables, 3 numerical variables and 1 Date variable; as seen in the figure 26.

```

> # Now has 128975 observations 18 columns
> # Datatype seems appropriate as well
> str(amazon_sale_rep_data)
'data.frame': 128975 obs. of 18 variables:
 $ Order.ID : chr "405-8078784-5731545" "171-9198151-1101146" "404-0687676-7273146" "403-9
615377-8133951" ...
 $ Date : Date, format: "2022-04-30" "2022-04-30" "2022-04-30" ...
 $ Status : chr "Cancelled" "Shipped - Delivered to Buyer" "Shipped" "Cancelled" ...
 $ Fulfilment : chr "Merchant" "Merchant" "Amazon" "Merchant" ...
 $ Sales.Channel : chr "Amazon.in" "Amazon.in" "Amazon.in" "Amazon.in" ...
 $ ship.service.level: chr "Standard" "Standard" "Expedited" "Standard" ...
 $ Style : chr "SET389" "JNE3781" "JNE3371" "J0341" ...
 $ SKU : chr "SET389-KR-NP-S" "JNE3781-KR-XXXL" "JNE3371-KR-XL" "J0341-DR-L" ...
 $ Category : chr "Set" "kurta" "kurta" "Western Dress" ...
 $ Size : chr "S" "3XL" "XL" "L" ...
 $ ASIN : chr "B09KXVBD7Z" "B09K3WFS32" "B07WV4JV4D" "B099NRCT7B" ...
 $ Courier.Status : chr NA "Shipped" "Shipped" NA ...
 $ Qty : int 0 1 0 1 1 1 0 1 ...
 $ Amount : num 648 406 329 753 574 ...
 $ ship.city : chr "MUMBAI" "BENGALURU" "NAVI MUMBAI" "PUDUCHERRY" ...
 $ ship.state : chr "MAHARASHTRA" "KARNATAKA" "MAHARASHTRA" "PUDUCHERRY" ...
 $ ship.postal.code : num 400081 560085 410210 605008 600073 ...
 $ B2B : chr "False" "False" "True" "False" ...

```

Figure 26

However, several duplicate values were found while performing analysis. Thus, I use the `distinct()` function to produce a dataset with only unique values. Therefore, now I have a dataset with 128969 observations, where 6 were duplicate values.

```

> # Remove duplicate values (6 of them)
> amazon_sale_rep_data <- amazon_sale_rep_data %>% distinct()
> str(amazon_sale_rep_data)
'data.frame': 128969 obs. of 18 variables:
 $ Order.ID : chr "405-8078784-5731545" "171-9198151-1101146" "404-0687676-7273146" "403-9
615377-8133951" ...
 $ Date : Date, format: "2022-04-30" "2022-04-30" "2022-04-30" ...
 $ Status : chr "Cancelled" "Shipped - Delivered to Buyer" "Shipped" "Cancelled" ...
 $ Fulfilment : chr "Merchant" "Merchant" "Amazon" "Merchant" ...
 $ Sales.Channel : chr "Amazon.in" "Amazon.in" "Amazon.in" "Amazon.in" ...
 $ ship.service.level: chr "Standard" "Standard" "Expedited" "Standard" ...
 $ Style : chr "SET389" "JNE3781" "JNE3371" "J0341" ...
 $ SKU : chr "SET389-KR-NP-S" "JNE3781-KR-XXXL" "JNE3371-KR-XL" "J0341-DR-L" ...
 $ Category : chr "Set" "kurta" "kurta" "Western Dress" ...
 $ Size : chr "S" "3XL" "XL" "L" ...
 $ ASIN : chr "B09KXVBD7Z" "B09K3WFS32" "B07WV4JV4D" "B099NRCT7B" ...
 $ Courier.Status : chr NA "Shipped" "Shipped" NA ...
 $ Qty : int 0 1 0 1 1 1 0 1 ...
 $ Amount : num 648 406 329 753 574 ...
 $ ship.city : chr "MUMBAI" "BENGALURU" "NAVI MUMBAI" "PUDUCHERRY" ...
 $ ship.state : chr "MAHARASHTRA" "KARNATAKA" "MAHARASHTRA" "PUDUCHERRY" ...
 $ ship.postal.code : num 400081 560085 410210 605008 600073 ...
 $ B2B : chr "False" "False" "True" "False" ...

```

Figure 27

I remove the missing values that were found from the dataset initially to avoid errors, anomalies or biases formed from missing values. Removing missing values also ensures that the complexity of the dataset is reduced.

```

> amazon_sale_rep_data <- na.omit(amazon_sale_rep_data)
> nrow(amazon_sale_rep_data)
[1] 116013

```

Figure 28

Hence, the summary and descriptive statistics of the variables in cleaned data are shown in Figure 29 below.

```

> # Summary statistics of cleaned data
> summary(amazon_cleaned)
  Order.ID      Date      Status      Fulfilment
Length:116013  Min. :2022-03-31  Length:116013  Length:116013
Class :character 1st Qu.:2022-04-20  Class :character  Class :character
Mode :character  Median :2022-05-10  Mode :character  Mode :character
                  Mean :2022-05-12
                  3rd Qu.:2022-06-05
                  Max. :2022-06-29
Sales.Channel  ship.service.level  Style      SKU      Category
Length:116013  Length:116013      Length:116013  Length:116013  Length:116013
Class :character  Class :character  Class :character  Class :character  Class :character
Mode :character  Mode :character  Mode :character  Mode :character  Mode :character

  Size      ASIN      Courier.Status      Qty      Amount
Length:116013  Length:116013  Length:116013  Min. :1.000  Min. : 0.0
Class :character  Class :character  Class :character  1st Qu.:1.000  1st Qu.: 449.0
Mode :character  Mode :character  Mode :character  Median :1.000  Median : 606.0
                  Mean :1.004  Mean : 649.8
                  3rd Qu.:1.000  3rd Qu.: 788.0
                  Max. :8.000  Max. :5584.0

  ship.city      ship.state      ship.postal.code      B2B
Length:116013  Length:116013  Min. :110001  Length:116013
Class :character  Class :character  1st Qu.:382424  Class :character
Mode :character  Mode :character  Median :500032  Mode :character
                  Mean :463320
                  3rd Qu.:600017
                  Max. :855117

```

Figure 29

Continuing the process of analysing our dataset, I manipulate the data by creating a revenue column in the dataset using the `mutate()` function. As mentioned in the introduction of this problem statement, the revenue is calculated by obtaining the product quantity and the amount of the item as shown in Figure 30.

```

> #creating revenue column
> amazon_cleaned <- amazon_cleaned %>%
+   mutate(Revenue = Qty * Amount)

```

Figure 30

Adding to that, I group the dataset based on the different categories of the product using the `group_by` function and output the total revenue using the `sum()` function on the newly created revenue column. Since I was unsure of the exact different categories and their product count in the dataset, I performed the `count()` function to output the number of products for each category. Hence, I realised that there are 9 different product categories as in Figure 31.

```

> #total revenue for each category
> total_revenue <- amazon_cleaned %>%
+   group_by(Category) %>%
+   summarise(Total_Revenue = sum(Revenue))
> head(total_revenue)
# A tibble: 6 x 2
  Category      Total_Revenue
  <chr>      <dbl>
1 Blouse      441259
2 Bottom      142870
3 Dupatta      915
4 Ethnic Dress 762949
5 Saree        125767
6 Set         37923384

> #counts in Category
> counts <- amazon_cleaned %>%
+   count(Category)
> print(counts)
  Category      n
1 Blouse      837
2 Bottom      393
3 Dupatta       3
4 Ethnic Dress 1050
5 Saree       148
6 Set      45077
7 Top       9864
8 Western Dress 13893
9 kurta    44748

```

Figure 31



To compare the revenue based on the number of products based on each category, we create a data frame while merging the total revenue and the counts using the Category column.

```
> #combine total revenue and counts
> merged_df <- merge(total_revenue, counts, by = "Category")
> head(merged_df)
```

	Category	Total_Revenue	n
1	Blouse	441259	837
2	Bottom	142870	393
3	Dupatta	915	3
4	Ethnic Dress	762949	1050
5	kurta	20667948	44748
6	Saree	125767	148

Figure 22

The average revenue per product was calculated by dividing the total\_revenue by the number of products sold. A data frame was created with a new column for the average revenue per product using the mutate() function.

```
> #finding the average revenue
> merged_df <- merged_df %>%
+ mutate(Average_Revenue_Per_Product = Total_Revenue / `Number of products sold`)
> head(merged_df)
```

	Category	Total_Revenue	Number of products sold	Average_Revenue_Per_Product
1	Blouse	441259	837	527.1912
2	Bottom	142870	393	363.5369
3	Dupatta	915	3	305.0000
4	Ethnic Dress	762949	1050	726.6181
5	kurta	20667948	44748	461.8742
6	Saree	125767	148	849.7770

Figure 33

## Data Visualisations

```
melt_bar <- melt(merged_df[,c('Category', 'Total_Revenue',
                              'Number of products sold')], id.vars = 1)
View(melt_bar)
merged_df2 <- merged_df %>%
  select(Category, Average_Revenue_Per_Product)
melt_bar2 <- left_join(melt_bar, merged_df2, by = "Category")
View(melt_bar2)
```


```
ggplot(data=melt_bar2, aes(x= reorder(Category, value), y=value, fill=variable, group=variable)) +
  geom_bar(position = 'dodge', stat='identity') +
  geom_text(aes(label=value), position=position_dodge(width=0.9), vjust=-0.25, size = 2.5) +
  scale_fill_manual(values = c("pink", "lightblue")) +
  scale_color_manual(values=c("blue")) +
  scale_y_log10() +
  geom_line(aes(x= reorder(Category, value), y = Average_Revenue_Per_Product,
                color = "Average_Revenue_Per_Product"), group = 1) +
  theme_bw() +
  labs(title = "Total Revenue vs Quantity Sold for each Product",
       x = "Products",
       y = "Frequency") +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=0.5),
        axis.text.y = element_blank(),
        plot.title = element_text(hjust = 0.5, size = 15),
        panel.grid = element_line(color = "white"),
        panel.background = element_rect(fill = "#ebf2ff"),
        legend.position = c(0.16, 0.88),
        legend.background = element_rect(fill = "#ebf2ff"),
        legend.text = element_text(size = 8),
        legend.title = element_blank(),
        legend.spacing.y = unit(-0.2, "cm"))
```

Figure 34

I reshaped the dataset using melt\_bar() to facilitate visualization, focusing on specific columns like Category, Total Revenue, and Number of products sold. Calculating the average revenue per product, I merged both datasets with left\_join(). The resulting side-by-side bar chart enabled straightforward category comparison within each group.



Employing ggplot, I constructed the chart: setting up initial parameters with `ggplot()`, adding adjacent bars for different variables with `geom_bar()`, and labelling values atop bars using `geom_text()`. I defined fill colours for bars with `scale_fill_manual()`, applied a logarithmic scale to the y-axis with `scale_y_log10()`, and plotted a line representing `Average_Revenue_Per_Product` with `geom_line()`. Customizations to plot appearance were made with `theme_bw()`, `labs()`, and `Theme()`.



*Figure 35*

This side-by-side bar chart illustrates both the total revenue and total quantity for each product sold, aiding the business in gauging demand and product popularity. Consequently, the business can implement tailored strategies such as price adjustments and targeted advertising campaigns. Additionally, the accompanying graph displays the average revenue per product, the contribution of each product to the overall revenue, for each category, calculated by dividing total revenue by quantity sold; this is important for informed decision- making.

### **Problem Statement 3 - Analysing the popularity of different colours across all products**

Analysing colour popularity across products is vital for inventory optimization and sales growth. This entails assessing colour popularity with stock levels and sales data to uncover trends and customer preferences. Detailed data on SKU codes, stock levels, sales quantities, and demand trends is essential. By identifying colours with high sales and consistent demand, businesses can optimize inventory, reduce holding costs, and improve sales. This analysis utilises the `amazon_sale_report.csv` and `sales_report.csv` datasets to enhance inventory turnover and meet customer demand more effectively.

#### **Data Cleaning & Manipulation**

The total observations in this dataset is 9271 and the total number of variables is 6. The sales report data contains 5 categorical variables and 1 numerical variable based on the datatypes shown in Figure 36.

```
> # Remove index (redundant)
> sale_rep_data <- select(sale_rep_data, -index)
> # Lots of empty rows at the bottom, datatype looks fine
> # 9271 observations
> str(sale_rep_data)
'data.frame':  9271 obs. of  6 variables:
 $ SKU.Code   : chr  "AN201-RED-L" "AN201-RED-M" "AN201-RED-S" "AN201-RED-XL" ...
 $ Design.No.: chr  "AN201" "AN201" "AN201" "AN201" ...
 $ Stock      : num  5 5 3 6 3 11 3 16 8 14 ...
 $ Category   : chr  "AN : LEGGINGS" "AN : LEGGINGS" "AN : LEGGINGS" "AN : LEGGINGS" ...
 $ Size       : chr  "L" "M" "S" "XL" ...
 $ Color      : chr  "Red" "Red" "Red" "Red" ...
```

*Figure 36*

There are no variables containing missing values (NAs). However, this also means that there are empty strings so there is a need to convert them to NA. After conversion and removing them using the `na.omit` function, the total number of observations after removing the missing values is 9188 as shown in Figure 37. I eliminate any missing values identified in the dataset initially to prevent errors, anomalies, or biases that may arise from them. This process also serves to simplify the dataset, reducing its complexity.

```

> # Not much NA detected, means they're empty spaces so convert them to NA
> sale_rep_data[sale_rep_data == ""] <- NA
> #remove NA values
> sale_rep_data <- na.omit(sale_rep_data)
> # 9188 observations now
> str(sale_rep_data)
'data.frame': 9188 obs. of 6 variables:
 $ SKU.Code : chr "AN201-RED-L" "AN201-RED-M" "AN201-RED-S" "AN201-RED-XL" ...
 $ Design.No.: chr "AN201" "AN201" "AN201" "AN201" ...
 $ Stock : num 5 5 3 6 3 11 3 16 8 14 ...
 $ Category : chr "AN : LEGGINGS" "AN : LEGGINGS" "AN : LEGGINGS" "AN : LEGGINGS" ...
 $ Size : chr "L" "M" "S" "XL" ...
 $ Color : chr "Red" "Red" "Red" "Red" ...
- attr(*, "na.action")= 'omit' Named int [1:83] 136 143 198 199 224 249 274 547 902 903 ...
..- attr(*, "names")= chr [1:83] "136" "143" "198" "199" ...

```

Figure 37

The summary and descriptive statistics of the variables in cleaned data are shown in Figure 38.

```

> sale_cleaned <- sale_rep_data
> # Summary statistics of cleaned data
> summary(sale_cleaned)

```

SKU.Code	Design.No.	Stock	Category	Size	Color
Length:9185	Length:9185	Min. : 0.00	Length:9185	Length:9185	Length:9185
Class :character	Class :character	1st Qu.: 3.00	Class :character	Class :character	Class :character
Mode :character	Mode :character	Median : 8.00	Mode :character	Mode :character	Mode :character
		Mean : 26.39			
		3rd Qu.: 31.00			
		Max. : 1234.00			

Figure 38

In view of analysis based on the problem statement, it was required to join the Amazon dataset together with the sales dataset to comment on the popularity of sales based on the product colours. I used both cleaned versions of the dataset to perform the left\_join function as seen in Figure 39. This function allows us to analyse only the products that have made sales.

```

> #PS3
> #rename SKU
> sale_cleaned <- sale_cleaned %>%
+   rename(SKU = SKU.Code)
> #join on amazon dataset and sales report
> combined_dataset <- left_join(amazon_cleaned, sale_cleaned, by = "SKU")

```

Figure 39

Since the analysis was on the sales made based on the colours of the product, I removed the unwanted columns from the combined datasets. Hence, this outputs the Colour and Category columns alone as seen below.

```

> #remove unwanted columns
> cleared_dataset <- combined_dataset[,c("Color", "Category.x")]
> cleared_dataset <- cleared_dataset %>%
+   filter(!is.na(Color))
> head(cleared_dataset)

```

	Color	Category.x
1	Green	kurta
2	Pink	kurta
3	Pink	kurta
4	Cream	Set
5	Light Green	kurta
6	Mauve	kurta

*Figure 40*

As seen in the categories, there are different types of products available in the dataset and it was necessary to find the number of products in each category with a specific colour. Thus, I performed a count function on each category based on each colour. The different categories include Kurta, Blouse, Bottom, Ethnic Dress, Saree, Set, Top and Western Dress. The figure below demonstrates the process of categorising the Kurta products based on the colour and their count. The same process was done for the other product categories.

```

> #most popular colors for kurta in descending order
> kurta_counts <- cleared_dataset %>%
+   filter(Category.x == "kurta") %>%
+   count(Color) %>%
+   arrange(desc(n)) %>%
+   rename(colour_count = n)
> head(kurta_counts)

```

	Color	colour_count
1	Blue	1892
2	Pink	964
3	Black	901
4	Green	747
5	Peach	629
6	Maroon	578

*Figure 41*

### Data Visualisations

A heatmap visually represents colour distribution across product categories, aiding in spotting trends and patterns. It uses a color-coded matrix to easily identify popular and less popular colours within each category. This helps determine if certain colours are consistently popular across multiple categories or if preferences vary by category.

```

# Create a list of all category counts with their respective category names
category_counts_list <- list(
  Kurta = kurta_counts,
  Blouse = blouse_counts,
  Bottom = bottom_counts,
  Ethnic Dress = ethnic_dress_counts,
  Saree = saree_counts,
  Set = set_counts,
  Top = top_counts,
  Western Dress = western_dress_counts
)

# Filter out NULL entries and combine the remaining data frames
all_counts <- lapply(names(category_counts_list), function(category) {
  df <- category_counts_list[[category]]
  if (is.null(df)) {
    df <- mutate(df, Category = category)
  }
  df
}) %>% bind_rows()

# Define a mapping from specific colors to broader color categories
color_mapping <- c(
  "Blue" = "Blue", "Teal" = "Blue", "Turquoise Blue" = "Blue", "Navy Blue" = "Blue",
  "Powder Blue" = "Blue", "Sky Blue" = "Blue", "Dark Blue" = "Blue", "Light Blue" = "Blue",
  "Navy" = "Blue", "Turquoise" = "Blue", "TEAL BLUE" = "Blue", "NAVY" = "Blue",
  "TEAL BLUE" = "Blue",
  "Pink" = "Pink", "Light Pink" = "Pink", "Magenta" = "Pink", "Hot Pink" = "Pink",
  "Coral Pink" = "Pink", "CORAL PINK" = "Pink",
  "Black" = "Black",
  "Green" = "Green", "Light Green" = "Green", "Dark Green" = "Green",
  "Olive Green" = "Green", "Sea Green" = "Green", "Olive" = "Green",
  "Turquoise Green" = "Green", "Teal Green" = "Green", "TEAL GREEN" = "Green",
  "Yellow" = "Yellow", "Gold" = "Yellow", "Light Yellow" = "Yellow",
  "LIGHT YELLOW" = "Yellow", "Mustard" = "Yellow", "Lemon" = "Yellow", "LEMON" = "Yellow",
  "LEMON" = "Yellow",
  "Red" = "Red", "Dark Red" = "Red", "Maroon" = "Red", "Wine" = "Red",
  "White" = "White", "OFF WHITE" = "White",
  "Grey" = "Grey", "Charcoal" = "Grey",
  "Peach" = "Peach", "Coral" = "Peach", "CORAL" = "Peach", "Coral Orange" = "Peach",
  "CORAL" = "Peach",
  "Orange" = "Orange", "ORANGE" = "Orange", "Rust" = "Orange", "CORAL ORANGE" = "Orange",
  "Purple" = "Purple", "Violet" = "Purple", "Mauve" = "Purple", "Indigo" = "Purple",
  "Brown" = "Brown", "Beige" = "Brown", "Tan" = "Brown", "Light Brown" = "Brown",
  "Multicolor" = "Multicolor",
  "Cream" = "Cream", "Taupe" = "Brown", "Lemon Yellow" = "Yellow", "Khaki" = "Brown",
  "LIME GREEN" = "Green", "BURGUNDY" = "Red", "MINT" = "Green", "MINT GREEN" = "Green"
)

```

Figure 42

After manipulating the dataset by creating a list of all category counts with their respective category names while making sure that the null data have been filtered out, I define a map from specific colours to broader categories.

```

# Map specific colors to broader color categories
all_counts$BroadColor <- color_mapping[all_counts$Color]

# Aggregate the data by BroadColor and Category
broad_color_counts <- all_counts %>%
  group_by(Category, BroadColor) %>%
  summarise(Count = sum(n, na.rm = TRUE)) %>%
  ungroup()

# Create the heatmap OF COLOUR CATEGORY
ggplot(broad_color_counts, aes(x = BroadColor, y = Category, fill = Count, label = Count)) +
  geom_tile() +
  geom_text(color = "black", size = 3) +
  scale_fill_gradient(low = "pink", high = "blue") +
  labs(title = "Heatmap of Color Popularity by Category",
       x = "Color",
       y = "Category") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1),
        plot.title = element_text(hjust = 0.5))

```

Figure 43

I used ggplot to create a heatmap showing colour popularity by category. The x-axis represents broad colours, the y-axis represents categories, and tile fill represents counts. I added text labels for counts and set a pink-to-blue gradient for fill colour. The plot's title is "Heatmap of Colour Popularity by Category", with x-axis labelled "Colour" and y-axis labelled "Category". I applied a minimal theme and customized it by rotating x-axis labels by 45 degrees and centring the plot title.




Figure 44

The heatmap depicts colour frequency across clothing types, illustrating Kurta as the most popular choice, indicated by the darkest shade of blue; which emerges as the top choice for tops and sets and the second most preferred for Western dresses, suggesting its overall popularity therefore the prevalence of blue might be skewed due to its various shades, with 13 categorized as different blues out of 71 total colours. This extensive categorization could inflate the perception of blues dominance compared to a more balanced categorization.

Bubble charts are visually appealing and simplify complex data by representing product categories as bubbles, with each bubbles size and colour indicating differences across categories and colours.

```
# Total counts for each broad color
total_color_counts <- all_counts %>%
  group_by(BroadColor) %>%
  summarise(TotalCount = sum(n, na.rm = TRUE)) %>%
  ungroup()

# Define the sizes corresponding to each broad color category
size_mapping <- c(1, 13, 6, 1, 9, 2, 1, 4, 5, 6, 4, 5, 2, 8)

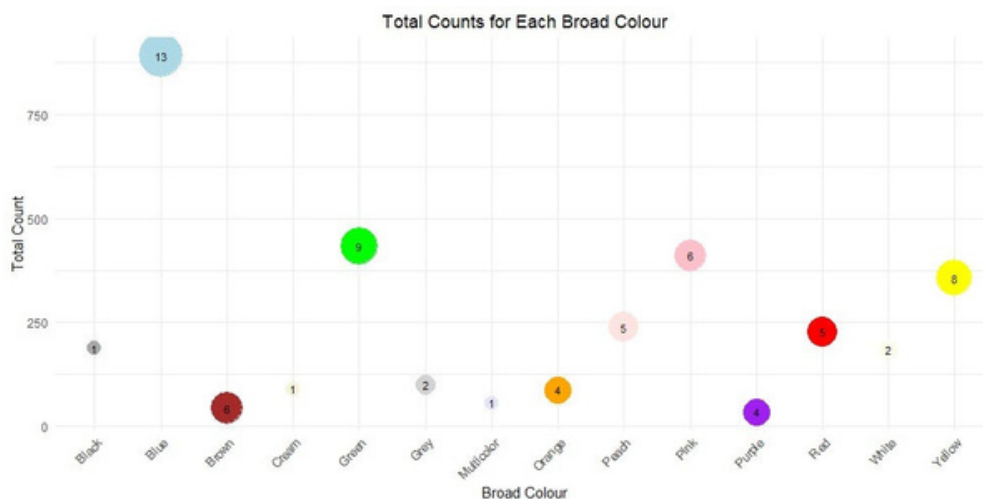
# Add a new column with the corresponding sizes
total_color_counts$BubbleSize <- size_mapping

# Create the bubble plot with customized aesthetics
ggplot(total_color_counts, aes(x = BroadColor, y = TotalCount, size = BubbleSize)) +
  geom_point(aes(color = BroadColor)) +
  geom_text(aes(label = BubbleSize, size = 0.8)) +
  labs(title = "Total Counts for Each Broad Colour",
       x = "Broad Colour",
       y = "Total Count",
       size = "Colour Categorisation Count") + # Add size legend label
  scale_size_continuous(range = c(3, 15), guide = guide_legend(title = "Colour Categorisation Count")) +
  scale_color_manual(values = c("darkgrey", "lightblue", "brown", "beige", "green", "lightgrey", "lavender",
                                "orange", "MistyRose", "pink", "purple", "red", "Ivory", "yellow")) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1), plot.title = element_text(hjust = 0.5)) +
  theme(legend.position = "none")
```

Figure 45



I began by aggregating the dataset using `group_by()` based on the "BroadColor" column, calculating the total count of each colour with `summarise()` while disregarding missing values. After ungrouping, each row represented a colour and its count. For the bubble chart, I customized bubble sizes by creating a vector and assigning them to a new column, "Bubble size", in the data frame. Using `ggplot`, I plotted the chart: points represented colours, sized by "Bubble size", coloured by "BroadColor", with text overlays showing bubble sizes. I labelled the plot using `labs()`, defined the size range with `scale_size_continuous()`, and assigned colours with `scale_color_manual()`. Applying a minimal theme with `theme_minimal()`, I further adjusted the appearance, including formatting x-axis text and centring the plot title. Lastly, I removed the legend with `theme(legend.position = "none")`.



*Figure 46*

This bubble chart illustrates the most frequent colours, with bubble height representing the frequency and bubble size indicating the number of items in each colour category (e.g., both navy blue and light blue fall under the blue category). From the analysis, a trend emerges: the higher the bubble, the larger the number within it that suggests that the way colours are categorized significantly influences the frequency shown; alternatively, it could be indicative of higher production due to more popularity amongst customers. Thus, a higher frequency for a colour group means there is more variation within that group. For instance, there are more types of blue shades because blue is a popular colour overall.

In conclusion, the company should prioritize improving customer retention rates to boost profits. This includes better inventory management and increased sales during peak months like October and February. Understanding product popularity informs pricing and advertising strategies. Focusing on producing clothing in popular colours like blue, with a variety of shades, can enhance sales.